

Modeling, Learning and Reasoning about Preference Trees over Combinatorial Domains

Xudong Liu

Advisor: Dr. Mirosław Truszczyński

Department of Computer Science
College of Engineering
University of Kentucky
Lexington, KY, USA
Tuesday, 4/19/2016

Preferences Are Ubiquitous

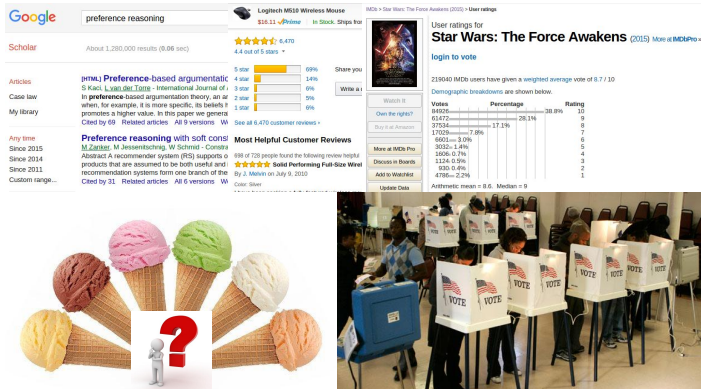


Figure : Preferences of different forms

Describing Preferences

Car1



<mvan, 7m, gray, big, honda, med, med>



Car2



<sedan, 5, blue, med, vw, med, med>

Figure : How to express preferences?

① How will I rate cars?

- For BodyType, I will assign 7 points to minivans, 5 to sedans, ...
- For Color, I will assign 8 points to blue, 4 to gray, ...

② What are the desired properties I see in cars?

- I prefer minivans to sedans, ...
- If minivan, I prefer gray to blue; if sedan, I prefer blue to gray; ...

Describing Preferences

Car1



<mvan, 7m, gray, big, honda, med, med>

Car2



<sedan, 5, blue, med, vw, med, med>



Figure : How to express preferences?

① How will I rate cars? (**Quantitative**)

- For BodyType, I will assign 7 points to minivans, 5 to sedans, ...
- For Color, I will assign 8 points to blue, 4 to gray, ...

② What are the desired properties I see in cars? (**Qualitative**)

- I prefer minivans to sedans, ...
- If minivan, I prefer gray to blue; if sedan, I prefer blue to gray; ...

Combinatorial Domains

Let \mathcal{I} be a finite set of attributes $\{X_1, \dots, X_p\}$, associated with a set of finite domains $\{Dom(X_1), \dots, Dom(X_p)\}$ for each attribute X_i . A *combinatorial domain* $CD(\mathcal{I})$ is a set of *objects* described by combinations of values from $Dom(X_i)$:

$$CD(\mathcal{I}) = \prod_{X_i \in \mathcal{I}} Dom(X_i).$$

Combinatorial Domains: Example

Domain of cars over set \mathcal{I} of p binary attributes:

① **BodyType**: {mvan, sedan}.

② **Capacity**: {5, 7m}.

③ **Color**: {blue, gray}.

⋮

$$CD(\mathcal{I}) = \underbrace{\{\langle \text{sedan}, 5, \text{blue}, \dots \rangle, \langle \text{mvan}, 7\text{m}, \text{gray}, \dots \rangle, \dots \}}_{2^p \text{ objects, too many!}}.$$

Combinatorial Domains: Example

Domain of cars:

- ① **BodyType**: {mvan, sedan, sport, suv}.
- ② **Capacity**: {2, 5, 7m}.
- ③ **Color**: {black, blue, gray, red, white}.
- ④ **LuggageSize**: {big, med, small}.
- ⑤ **Make**: {bmw, ford, honda, vw}.
- ⑥ **Price**: {low, med, high, vhigh}.
- ⑦ **Safety**: {low, med, high}.

Single Agent

Car1



<mvan, 7m, gray, big, honda, med, med>

Car2



<sedan, 5, blue, med, vw, med, med>



Figure : Dominance and Optimization

Multi-Agent



Figure : Social Choice and Welfare

Research Problems of Interest

- ① Preference representation formalisms to compactly model qualitative preferences over combinatorial domains.
- ② Preference elicitation and learning methods to cast preferences of agents as a theory in a preference formalism.
- ③ Preference reasoning tasks:
 - Dominance and optimization
 - Preference aggregation

Q: How do we compactly represent qualitative preferences over combinatorial domains?

- ① Preference Trees (P-trees)^{1,11}
- ② Partial Lexicographic Preference Trees (PLP-trees)⁸
- ③ Lexicographic Preference Trees (LP-trees)^{4,13}

¹Niall M Fraser. "Ordinal preference representations". In: Theory and Decision (1994)

²Xudong Liu and Mirosław Trzuszczynski. "Preference Trees: A Language for Representing and Reasoning about Qualitative Preferences". In: Proceedings of the 8th Multidisciplinary Workshop on Advances in Preference Handling (MPREF). 2014

³Xudong Liu and Mirosław Trzuszczynski. "Learning Partial Lexicographic Preference Trees over Combinatorial Domains". In: Proceedings of the 29th AAAI Conference on Artificial Intelligence (AAAI). 2015

⁴Richard Booth et al. "Learning conditionally lexicographic preference relations". In: ECAI. 2010

⁵Xudong Liu and Mirosław Trzuszczynski. "Aggregating Conditionally Lexicographic Preferences Using Answer Set Programming Solvers". In: Proceedings of the 3rd International Conference on Algorithmic Decision Theory (ADT). 2013

Q: How do we learn predictive qualitative preference models over combinatorial domains?

- ① Partial Lexicographic Preference Trees (PLP-trees)^{6,7,8}
 - Active and passive learning
 - Compute a (possibly small) PLP-tree consistent with all the data
 - Compute a PLP-tree that agrees with the data as much as possible
- ② Empirical Learning of PLP-trees and PLP-forests⁹

⁶Michael Schmitt and Laura Martignon. "On the complexity of learning lexicographic strategies". In: The Journal of Machine Learning Research (2006)

⁷József Dombi, Csanád Imreh, and Nándor Vincze. "Learning lexicographic orders". In: European Journal of Operational Research (2007)

⁸Xudong Liu and Mirosław Truszczynski. "Learning Partial Lexicographic Preference Trees over Combinatorial Domains". In: Proceedings of the 29th AAAI Conference on Artificial Intelligence (AAAI). 2015

⁹Xudong Liu and Mirosław Truszczynski. "Learning Partial Lexicographic Preference Trees and Forests over Multi-Valued Attributes". In: Review by ECAI-16 Program Committee

Q: How do we reason about preferences over combinatorial domains?

① Preference Optimization^{10,11,12}:

- Dominance testing: $o_1 \succ_P o_2$?
- Optimality testing: $o_1 \succ_P o_2$ for all $o_2 \neq o_1$?
- Optimality computing: what is the optimal outcome wrt P ?

② Preference Aggregation¹³:

- Winner determination: which candidate wins the election?
- “Strong” candidate: a candidate with score more than a threshold?

¹⁰Jérôme Lang, Jérôme Mengin, and Lirong Xia. “Aggregating Conditionally Lexicographic Preferences on Multi-issue Domains”. In: CP. 2012

¹¹Xudong Liu and Mirosław Trzuszczynski. “Preference Trees: A Language for Representing and Reasoning about Qualitative Preferences”. In: Proceedings of the 8th Multidisciplinary Workshop on Advances in Preference Handling (MPREF). 2014

¹²Xudong Liu and Mirosław Trzuszczynski. “Reasoning with Preference Trees over Combinatorial Domains”. In: Proceedings of the 4th International Conference on Algorithmic Decision Theory (ADT). 2015

¹³Xudong Liu and Mirosław Trzuszczynski. “Aggregating Conditionally Lexicographic Preferences Using Answer Set Programming Solvers”. In: Proceedings of the 3rd International Conference on Algorithmic Decision Theory (ADT). 2013

- ① Modeling qualitative preferences:
 - Preference trees (P-trees)
 - Partial lexicographic preference trees (PLP-trees)
- ② Learning PLP-trees and PLP-forests
- ③ Aggregating LP-trees
- ④ Future research directions

- ① Modeling qualitative preferences:
 - Preference trees (P-trees)
 - Partial lexicographic preference trees (PLP-trees)
- Learning PLP-trees and PLP-forests
- Aggregating LP-trees
- Future research directions

Preference Trees

- 1 Let $\mathcal{I} = \{X_1, \dots, X_p\}$ be a set of attributes, and $D(\mathcal{I}) = \{Dom(X_1), \dots, Dom(X_p)\}$ a set of finite domains for \mathcal{I} .
- 2 A *literal* is an assignment to an attribute. We denote by $X_i := x_{i,j}$ the literal that assigns value $x_{i,j} \in Dom(X_i)$ to X_i . When no confusion, we write $x_{i,j}$, instead of $X_i := x_{i,j}$, as a literal. We then denote by $\mathcal{L} = \{x_{i,j} \in Dom(X_i) : X_i \in \mathcal{I}\}$ the set of literals given \mathcal{I} and $D(\mathcal{I})$.
- 3 The combinatorial domain $CD(\mathcal{I})$ is defined as earlier.

- 4 A **P-tree** T over $CD(\mathcal{I})$ is a binary tree, where non-leaf nodes are labeled with propositional formulas over \mathcal{L} .
- 5 Given an outcome $o \in CD(\mathcal{I})$, the **leaf** $l_T(o)$ is the leaf reached by traversing the tree T according to o . When at a node N labeled with φ , if $o \models \varphi$, we descend to the left child of N ; otherwise, to the right.
- 6 For $o_1, o_2 \in CD(\mathcal{I})$, we have $o_1 \succ_T o_2$ if $l_T(o_1) \succ_T l_T(o_2)$, and $o_1 \approx_T o_2$ if $l_T(o_1) = l_T(o_2)$. Outcome o_1 is **optimal** if there exists no o_2 such that $o_2 \succ_T o_1$.

Preference Trees (P-Trees)

Let φ , ψ , and π be propositional formulas over the set \mathcal{L} of literals that are values from $\bigcup_{X_i \in V} \text{Dom}(X_i)$.




Figure : A P-tree

$$\varphi \wedge \psi \succ \varphi \wedge \neg\psi \succ \neg\varphi \wedge \pi \succ \neg\varphi \wedge \neg\pi.$$

Preference Trees (P-Trees)

Let φ , ψ , and π be propositional formulas over the set \mathcal{L} of literals that are values from $\bigcup_{X_i \in V} \text{Dom}(X_i)$.




Figure : A P-tree

$$\varphi \wedge \psi \succ \varphi \wedge \neg\psi \succ \neg\varphi \wedge \pi \succ \neg\varphi \wedge \neg\pi.$$

Total preorder

Example: The Cars Domain

- ① **BodyType**(X_1): {mvan($x_{1,1}$), sedan($x_{1,2}$), sport($x_{1,3}$), suv($x_{1,4}$)}.
- ② **Capacity**(X_2): {2, 5, 7m}.
- ③ **Color**(X_3): {black, blue, gray, red, white}.
- ④ **LuggageSize**(X_4): {big, med, small}.
- ⑤ **Make**(X_5): {bmw, ford, honda, vw}.
- ⑥ **Price**(X_6): {low, med, high, vhigh}.
- ⑦ **Safety**(X_7): {low, med, high}.

Example: Preference Trees over Cars

BodyType(X_1): {mvan($x_{1,1}$), sedan($x_{1,2}$), sport($x_{1,3}$), suv($x_{1,4}$)}.

Color(X_3): {black, blue, gray, red, white}.

Price(X_6): {low, med, high, vhigh}.




Figure : A P-tree over cars¹⁴

¹⁴ $\varphi = (x_{1,1} \wedge x_{3,5}) \vee (x_{1,2} \wedge x_{3,2})$.

Example: Preference Trees over Cars

BodyType(X_1): {mvan($x_{1,1}$), sedan($x_{1,2}$), sport($x_{1,3}$), suv($x_{1,4}$)}.

Color(X_3): {black, blue, gray, red, white}.

Price(X_6): {low, med, high, vhigh}.




Figure : A P-tree over cars¹⁴

Car2 \succ *Car1*


¹⁴ $\varnothing = (x_{1,1} \wedge x_{3,5}) \vee (x_{1,2} \wedge x_{3,2})$.

Compact Representation of P-trees


BodyType(X_1): {mvan($x_{1,1}$), sedan($x_{1,2}$), sport($x_{1,3}$), suv($x_{1,4}$)}.

Color(X_3): {black, blue, gray, red, white}.

Price(X_6): {low, med, high, vhigh}.




(a) Full




(b) Compact

Figure : Compact P-trees

Compact Representation of P-trees



(a) Full



(b) Compact

Figure : Compact P-trees

Compact Representation of P-trees

A *compact P-tree* over $CD(\mathcal{I})$ is a binary tree where

- 1 every node is labeled with a Boolean formula over \mathcal{I} , and
- 2 every non-leaf node t labeled with φ has either two outgoing edges (Fig. (a)), or one outgoing edge pointing straight-down (Fig. (b)), left (Fig. (c)), or right (Fig. (d)).







Figure : Compact P-trees

Relative Expressivity of Preference Languages



Computational Complexity Results

Dominance-testing (DOMTEST): $o_1 \succ_T o_2$?

Optimality-testing (OPTTEST): o optimal w.r.t T ?

Optimality-with-property (OPTPROP): is there optimal o with property α ?

- ① DOMTEST $\in P$
- ② OPTTEST $\in coNP$ -complete:
 - The complement problem is reduced from the SAT problem.
- ③ OPTPROP $\in \Delta_2^P$ -complete:
 - The problem is reduced from the Maximum Satisfying Assignment (MSA) problem.

- ① Modeling qualitative preferences:
 - Preference trees (P-trees)
 - Partial lexicographic preference trees (PLP-trees)
- Learning PLP-trees and PLP-forests
- Aggregating LP-trees
- Future research directions

The Cars Domain

- ① **BodyType(B)**: {mvan, sedan, sport, suv}.
- ② **Capacity(C)**: {2, 5, 7m}.
- ③ **Color(O)**: {black, blue, gray, red, white}.
- ④ **LuggageSize(L)**: {big, med, small}.
- ⑤ **Make(M)**: {bmw, ford, honda, vw}.
- ⑥ **Price(P)**: {low, med, high, vhigh}.
- ⑦ **Safety(S)**: {low, med, high}.

Partial Lexicographic Preference Trees (PLP-Trees)

A *PLP-tree* over $CD(\mathcal{I})$ is a tree, where

- 1 every non-leaf node t is labeled with an attribute $Attr(t)$ in \mathcal{I} ,
- 2 every non-leaf node t has $|Dom(Attr(t))|$ outgoing edges labeled with a value of $Attr(t)$, and
- 3 every attribute appears *at most* once on every branch.




Figure : A PLP-tree over cars

Partial Lexicographic Preference Trees (PLP-Trees)

A *PLP-tree* over $CD(\mathcal{I})$ is a tree, where

- 1 every non-leaf node t is labeled with an attribute $Attr(t)$ in \mathcal{I} ,
- 2 every non-leaf node t has $|Dom(Attr(t))|$ outgoing edges labeled with a value of $Attr(t)$, and
- 3 every attribute appears *at most* once on every branch.




Figure : A PLP-tree over cars

Total preorder

Partial Lexicographic Preference Trees (PLP-Trees)

Car1



<mvan, 7m, gray, big, honda, med, med>

Car2



<sedan, 5, blue, med, vw, med, med>






Figure : A PLP-tree over cars

Car2 \succ Car1

Compact Representations of PLP-Tree




(a) Full




(b) Compact

Figure : Unconditional Importance & Unconditional Preference (UIUP)

Compact Representations of PLP-Tree



(a) Full



(b) Compact

Figure : Unconditional Importance & Conditional Preference (UICP)

Compact Representations of PLP-Tree




Figure : Conditional Importance & Unconditional Preference (CIUP)

Lexicographic Preference Trees (LP-Trees)

- ① An *LP-tree* \mathcal{L} over $CD(\mathcal{I})$ is a PLP-tree, where
- each attribute appears **exactly once** on every path from the root to a leaf.
 - Unlike PLP-trees, an LP-tree induces a total order.

Conclusion

- ① Generalizing LP-trees, PLP-trees compactly represent total *preorders* over combinatorial domains, by allowing agents to specify, on each path, only a subset of attributes (i.e., those useful ones).
- ② P-trees further generalize PLP-trees by labeling the nodes with *propositional formulas*, in practice, usually built with small number (e.g., at most 3) of attributes.


- Modeling qualitative preferences:
 - Preference trees (P-trees)
 - Partial lexicographic preference trees (PLP-trees)
- ② Learning PLP-trees and PLP-forests
- Aggregating LP-trees
- Future research directions

Learning PLP-trees

Consistent Learning (CONSLearn)

Given an example set \mathcal{E} , decide whether there exists a PLP-tree T (of a particular type) such that T is consistent with \mathcal{E} .

($\langle \text{sdn}, 5, \text{blk}, m, h, m, m \rangle, \langle \text{suv}, 7m, \text{wht}, b, f, m, m \rangle$)
($\langle \text{spt}, 2, \text{wht}, s, \text{bmw}, h, h \rangle, \langle \text{spt}, 2, \text{wht}, s, \text{bmw}, \text{vh}, h \rangle$)
($\langle \text{mvn}, 7m, \text{gry}, b, f, m, m \rangle, \langle \text{sdn}, 5, \text{bl}, m, f, m, m \rangle$)



UIUP tree

Learning PLP-trees

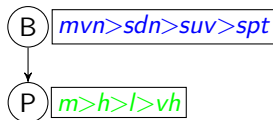
Small Learning (SMALLLEARN)

Given an example set \mathcal{E} and a positive integer l ($l \leq |\mathcal{E}|$), decide whether there exists a PLP-tree T (of a particular type) such that T is consistent with \mathcal{E} and $|T| \leq l$.

($\langle \text{sdn}, 5, \text{blk}, m, h, m, m \rangle, \langle \text{suv}, 7m, \text{wht}, b, f, m, m \rangle$)

($\langle \text{spt}, 2, \text{wht}, s, \text{bmw}, h, h \rangle, \langle \text{spt}, 2, \text{wht}, s, \text{bmw}, \text{vh}, h \rangle$)

($\langle \text{mvn}, 7m, \text{gry}, b, f, m, m \rangle, \langle \text{sdn}, 5, \text{bl}, m, f, m, m \rangle$)



UIUP tree

Learning PLP-trees

Maixmal Learning (MAXLEARN)


Given an example set \mathcal{E} and a positive integer k ($k \leq m$), decide whether there exists a PLP-tree T (of a particular type) such that T satisfies at least k examples in \mathcal{E} .

($\langle \text{sdn}, 5, \text{blk}, m, h, m, m \rangle, \langle \text{suv}, 7m, \text{wht}, b, f, m, m \rangle$)

($\langle \text{spt}, 2, \text{wht}, s, \text{bmw}, h, h \rangle, \langle \text{spt}, 2, \text{wht}, s, \text{bmw}, \text{vh}, h \rangle$)

($\langle \text{mvn}, 7m, \text{gry}, b, f, m, m \rangle, \langle \text{sdn}, 5, \text{bl}, m, f, m, m \rangle$)

($\langle \text{suv}, 7m, \text{gry}, b, \text{vw}, \text{vh}, m \rangle, \langle \text{suv}, 7m, \text{gry}, b, \text{vw}, h, m \rangle$)



UIUP tree

Complexity Results on PLP-trees

	UP	CP
UI	P	NP
CI	NPC ¹⁵	P

(a) CONSLearn

	UP	CP
UI	NPC	NPC
CI	NPC	NPC

(b) SMALLLearn

	UP	CP
UI	NPC ¹⁶	NPC
CI	NPC	NPC

(c) MAXLearn

Figure : Complexity results for learning PLP-trees

¹⁵Booth et al., *Learning Conditionally Lexicographic Preference Relations*, 2010.

¹⁶Schmitt and Martignon, *On the Complexity of Learning Lexicographic Strategies*, 2006.

Experimentation

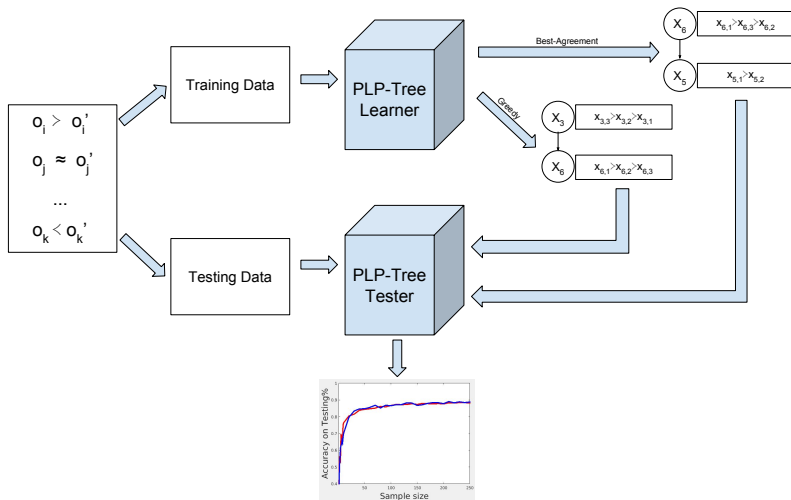


Figure : PLP-tree learning system

Datasets

Dataset	p	$ \mathcal{X} $	$ \mathcal{E}^> $	$ \mathcal{E}^{\approx} $
BreastCancerWisconsin	9	270	9,009	27,306
CarEvaluation	6	1,728	682,721	809,407
CreditApproval	10	520	66,079	68,861
GermanCredit	10	914	172,368	244,873
Ionosphere	10	118	3,472	3,431
MammographicMass	5	62	792	1,099
Mushroom	10	184	8,448	8,388
Nursery	8	1,266	548,064	252,681
SPECTHeart	10	115	3,196	3,359
TicTacToe	9	958	207,832	250,571
Vehicle	10	455	76,713	26,572
Wine	10	177	10,322	5,254

Figure : Preference Learning Library¹⁷


¹⁷<http://www.cs.uky.edu/~liu/preflearnlib.php>

Experimental Results: Best-Agreement vs. Greedy


Dataset	BA-UIUP	G-UIUP
BCW	88.4	88.2
CE	84.8	83.6
CA	91.1	89.3
GC	72.2	72.2
IN	87.0	79.6
MM	87.5	86.8
MS	84.8	70.3
NS	91.8	91.7
SH	93.2	92.6
TTT	72.1	71.9
VH	76.8	76.6
WN	96.0	95.5

Table : Accuracy (percentage of correctly handled testing examples) for UIUP PLP-trees learned using the best-agreement and the greedy methods on the learning data (250 of \mathcal{E}^{\sim})


Experimental Results: Best-Agreement vs. Greedy




(a) CarEvaluation



(b) Ionosphere



(c) Mushroom



(d) Wine

Figure : Learning curves solving MAXLEARN for UIUP PLP-trees

Experimental Results: Greedy

Dataset	UIUP	UICP-1	CIUPB	CIUPD	CICP
BCW	90.7	91.4	91.0	90.7	91.4
CE	85.8	86.0	85.8	85.9	86.0
CA	91.4	91.7	91.6	92.0	92.2
GC	74.3	74.6	74.3	74.5	75.7
IN	87.1	86.9	87.2	88.5	90.4
MM	88.2	89.5	87.3	86.9	90.0
MS	71.6	74.2	77.1	75.6	76.6
NS	92.9	93.0	93.0	93.0	93.0
SH	93.4	94.9	95.4	94.8	95.7
TTT	73.9	74.5	74.4	75.4	76.2
VH	79.2	80.4	80.3	80.0	81.2
WN	95.5	97.8	97.8	97.5	97.8

Table : Accuracy percents on the testing data (30% of \mathcal{E}^{\succ}) for all four classes of PLP-trees, using models learned by the greedy algorithm from the learning data (the other 70% of \mathcal{E}^{\succ})

Experimental Results: Sizes of PLP-trees by Greedy

Dataset	UIUP	UICP-1	CI	$ \mathcal{E}_{train}^> $
BCW	9	33	87,381	6,306
CE	6	21	853	477,904
CA	10	37	91,477	46,255
GC	10	37	349,525	120,657
IN	10	19	1,023	2,430
MM	5	17	341	554
MS	10	37	91,477	5,913
NS	8	29	7,765	383,644
SH	10	19	1,023	2,237
TTT	9	25	9,841	145,482
VH	10	37	349,525	53,699
WN	10	37	349,525	7,225

Table : Maximum sizes of trees for all the classes and the training sample sizes for all datasets

Experimental Results: Sizes of PLP-trees by Greedy

Dataset	UIUP	UICP-1	CIUPB	CIUPD	CICP
BCW	6.7	21.8	19.8	28.0	25.7
CE	6.0	17.0	73.2	108.9	109.5
CA	9.0	24.7	31.3	78.6	81.1
GC	9.7	36.0	49.8	210.3	190.0
IN	9.6	17.2	19.8	31.5	30.6
MM	4.5	14.7	8.3	10.8	10.0
MS	7.6	20.7	15.7	22.7	16.3
NS	8.0	25.7	56.2	121.0	116.9
SH	8.4	13.7	13.0	18.4	19.0
TTT	8.0	21.8	36.8	126.8	115.2
VH	9.0	32.7	33.9	101.3	105.4
WN	5.1	13.3	14.2	16.9	14.6

Table : Sizes of trees learned by the greedy algorithm from the training data (70% of \mathcal{E}^{\sim})

Partial Lexicographic Preference Forests (PLP-Forests)

- ① Inspired by *random forests*, we proposed *PLP-forests* that are sets of PLP-trees; thus, the four classes.
- ② To reduce the overfitting of a PLP-tree, a PLP-forest
 - consists of *diverse* trees (learned from *small* training samples), and
 - aggregates its constituent trees using the *Pairwise Majority Rule* (PMR).

Experimental Results: Best-Agreement vs. Greedy

Dataset	G+Tree	G+Forest	BA+Forest
BCW	90.7	93.4	95.1
CE	85.8	91.9	89.2
CA	91.4	91.5	93.1
GC	74.3	75.4	77.9
IN	87.1	83.0	92.5
MM	88.2	89.1	90.8
MS	71.6	78.8	90.2
NS	92.9	93.2	94.0
SH	93.4	93.7	94.9
TTT	73.9	75.1	77.2
VH	79.2	82.7	81.9
WN	95.5	95.8	96.9


Table : Accuracy percents on the testing data (30% of \mathcal{E}^{\succ}) for UIUP trees and forests of 5000 UIUP trees, using the greedy and the best-agreement algorithms from the learning data (the other 70% of \mathcal{E}^{\succ})

Experimental Results: Greedy


Dataset	UIUP	UICP-1	CIUPB	CIUPD	CICP
BCW	93.4	94.1	93.7	94.1	94.0
CE	91.9	88.3	91.4	89.7	91.4
CA	91.5	91.6	92.8	92.9	93.0
GC	75.4	73.8	76.1	76.1	76.2
IN	83.0	87.9	89.3	89.4	89.5
MM	89.1	90.1	90.0	90.1	90.2
MS	78.8	87.2	92.2	92.2	91.8
NS	93.2	89.9	93.3	93.4	93.4
SH	93.7	93.5	93.6	93.6	93.7
TTT	75.1	75.2	76.6	76.5	76.9
VH	82.7	81.8	83.2	83.2	83.4
WN	95.8	95.4	97.5	97.8	97.8

Table : Accuracy percents on the testing data (30% of \mathcal{E}^{\succ}) for all four classes of PLP-forests of 5000 trees, using the greedy algorithm from the learning data (the other 70% of \mathcal{E}^{\succ})


Experimental Results: UIUP vs. CICP




(a) CarEvaluation



(b) Ionosphere



(c) Mushroom



(d) Wine

Figure : Greedy learning curves solving MAXLEARN for PLP-forests

Conclusion

- ① PLP-trees and PLP-forests are *expressive* preference models.
- ② PLP-forests aggregated by PRM provide in general *higher* accuracy than PLP-trees.
- ③ PLP-trees and PLP-forests learned by a greedy approximation method have accuracy *comparable* to best-agreement PLP-trees and PLP-forests.
- ④ The greedy algorithms are *fast*, can work with *large* datasets (of \sim half million examples), and can compute *small* models.

- Modeling qualitative preferences:
 - Preference trees (P-trees)
 - Partial lexicographic preference trees (PLP-trees)
- Learning PLP-trees and PLP-forests
- ③ Aggregating LP-trees
- Future research directions

Positional Scoring Rules

- k -approval: $(1, \dots, 1, 0, \dots, 0)$ with k being the number of 1's.
- (k, l) -approval: $(c, \dots, c, d, \dots, d, 0, \dots, 0)$, where c and d are constants ($c > d$), and the numbers of c 's and d 's equal to k and l .
- b -Borda: $(b, b - 1, \dots, b - m + 1)$, where b is a constant and m is the number of candidates.

The Evaluation and Winner Problems

The Evaluation Problem

Let r be a positional scoring rule with a scoring vector w , \mathcal{C} a class of LP-trees. Given a \mathcal{C} -profile P of n LP-trees over p attributes and a positive integer R , the *evaluation* problem is to decide whether there exists an alternative $o \in \mathcal{X}$ such that $s_w(o, P) \geq R$.

The Winner Problem

Let r be a positional scoring rule with a scoring vector w , \mathcal{C} a class of LP-trees. Given a \mathcal{C} -profile P of n LP-trees over p attributes, the *winner* problem is to compute an alternative $o \in \mathcal{X}$ with the maximum score $s_w(o, P)$.

Complexity of the Evaluation Problem: k -Approval

	UP	CP
UI	P	P
CI	P	P

(a) $k = 2^{p-1} \pm f(p)$, $f(p)$ is a poly

	UP	CP
UI	NPC	NPC
CI	NPC	NPC

(b) $k = 2^{p-c}$, $c > 1$ is a const

Figure : k -Approval

Complexity of the Evaluation Problem: (k, l) -Approval

	UP	CP
UI	P	P
CI	P	P

(a) $k = l = 2^{p-1}$

	UP	CP
UI	NPC	NPC
CI	NPC	NPC

(b) $k = l = 2^{p-c}$, $c > 1$ is a const

Figure : (k, l) -Approval

Complexity of the Evaluation Problem: b -Borda

	UP	CP
UI	P	NPC
CI	NPC	NPC

(a) $b = 2^p - 1$

	UP	CP
UI	NPC	NPC
CI	NPC	NPC

(b) $b = 2^{p-c} - 1$, $c \geq 1$ is a const

Figure : b -Borda

Modeling the Problems in ASP

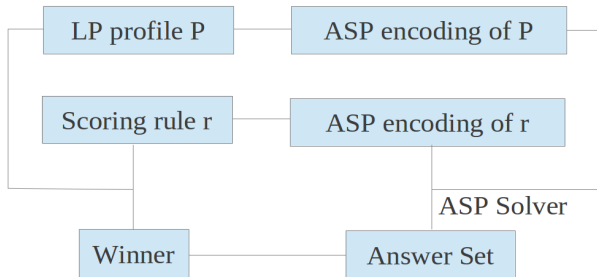


Figure : The winner problem

- Solvers: *clingo*¹⁸, *clingcon*¹⁹

¹⁸M. Gebser et al. "Potassco: The Potsdam Answer Set Solving Collection". In: AI Communications (2011)

¹⁹Max Ostrowski and Torsten Schaub. "ASP modulo CSP: The clingcon system". In: TPLP (2012)

Modeling the Problems in W-MAXSAT

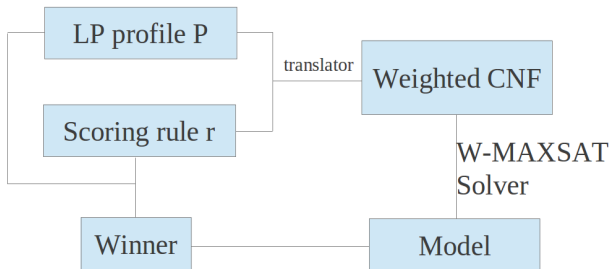



Figure : The winner problem


- Solver: *toulbar*²⁰

²⁰M Sanchez et al. "Max-CSP competition 2008: toulbar2 solver description". In: the Third International CSP Solver Competition (2008)

Varying p and n : 2^{p-2} -approval




(a) Fixed #attributes (10)




(b) Fixed #votes (1000)

Figure : Solving the winner problem

Varying p and n : $(2^{p-2}, 2^{p-2})$ -approval ²¹



(a) Fixed #attributes (10)




(b) Fixed #votes (1000)


Figure : Solving the winner problem

²¹ scoring vector: $(2, \dots, 2, 1, \dots, 1, 0, \dots, 0)$ with the numbers of 2's and 1's equal to 2^{p-2}

Varying p and n : Borda



(a) Fixed #attributes (10)



(b) Fixed #votes (1000)

Figure : Solving the winner problem

Conclusion

- ① When votes are total orders of candidates, computing a winner for a positional voting rule is computationally easy; not necessarily so, when they are LP-trees over combinatorial domains.
- ② For the cases when determining a winner is computationally hard, we solved this problem using ASP and empirically showed that ASP tools are *effective* on large instances.

- Modeling qualitative preferences:
 - Preference trees (P-trees)
 - Partial lexicographic preference trees (PLP-trees)
- Learning PLP-trees and PLP-forests
- Aggregating LP-trees
- ④ Future research directions

Data-Driven Preference Learning:

① Recommender Systems²²:

- Collaborative
- Content-based
- Hybrid

② Machine Learning (fitting function):

- Supervised learning (e.g., decision trees, random forests)
- Label ranking²³

③ Model-based Learning (learning interpretable decision models):

- Preference Elicitation (Human-in-the-Loop)
- Conditional Preference Networks, Preference Trees
- Stochastic Models (e.g., Choquet integral²⁴, TOPSIS-like models²⁵)

²²Gediminas Adomavicius and Alexander Tuzhilin. "Toward the next generation of recommender systems: A survey of the state-of-the-art and possible extensions". In: Knowledge and Data Engineering, IEEE Transactions on (2005)

²³Eyke Hüllermeier et al. "Label ranking by learning pairwise preferences". In: Artificial Intelligence (2008)

²⁴Ali Fallah Tehrani, Weiwei Cheng, and Eyke Hüllermeier. "Choquistic Regression: Generalizing Logistic Regression using the Choquet Integral." In: EUSFLAT. 2011

²⁵Manish Agarwal, Ali Fallah Tehrani, and Eyke Hüllermeier. "Preference-based Learning of Ideal Solutions in TOPSIS-like Decision Models". In: Journal of Multi-Criteria Decision Analysis (2014)

Preference Reasoning and Applications:

- ① Social Choice and Welfare^{26,27}:
 - Voting
 - Fair division
 - Strategyproof Social Choice
- ② Automated Planning and Scheduling^{28,29,30}:
 - Travel scheduling
 - Manufacturing
 - Traffic control

²⁶Kenneth J Arrow, Amartya Sen, and Kotaro Suzumura. Handbook of Social Choice and Welfare. Vol. 1 & 2. 2010

²⁷Felix Brandt, Vincent Conitzer, and Ulle Endriss. "Computational social choice". In: Multiagent systems (2012)

²⁸Tran Cao Son and Enrico Pontelli. "Planning with preferences using logic programming". In: Theory and Practice of Logic Programming (2006)

²⁹Meghyn Bienvenu, Christian Fritz, and Sheila A McIlraith. "Specifying and computing preferred plans". In: Artificial Intelligence (2011)

³⁰Hannah Bast et al. "Route planning in transportation networks". In: arXiv preprint (2015)

① Quantitative:

- Utility/Cost Functions³¹
- Possibilistic Logic³²
- Fuzzy Preference Relations³³
- Penalty Logic³⁴

② Qualitative:

- Answer-Set Optimization Theories³⁵
- Ceteris Paribus Networks (e.g., CP-nets³⁶, TCP-nets³⁷, CI-nets³⁸)
- Conditional Preference Theories³⁹

³¹Souhila Kaci. Working with Preferences: Less Is More: Less Is More. Springer Science & Business Media, 2011

³²Didier Dubois, Jérôme Lang, and Henri Prade. "A Brief Overview of Possibilistic Logic". In: ECSQARU. 1991

³³SA Orlovsky. "Decision-making with a fuzzy preference relation". In: Fuzzy sets and systems (1978)

³⁴Gadi Pinkas. Propositional non-monotonic reasoning and inconsistency in symmetric neural networks. 1991

³⁵Gerhard Brewka, Ilkka Niemelä, and Mirosław Truszczyński. "Answer Set Optimization". In: IJCAI. 2003

³⁶C. Boutilier et al. "CP-nets: A Tool for Representing and Reasoning with Conditional Ceteris Paribus Preference Statements". In: Journal of Artificial Intelligence Research (2004)

³⁷Ronen I. Brafman and Carmel Domshlak. "Introducing Variable Importance Tradeoffs into CP-Nets". In: UAI. 2002

³⁸Sylvain Bouveret, Ulle Endriss, and Jérôme Lang. "Conditional importance networks: A graphical language for representing ordinal, monotonic preferences over sets of goods". In: (2009)

³⁹Nic Wilson. "Extending CP-Nets with Stronger Conditional Preference Statements". In: AAAI-04. 2004

Summary

- 1 The language of P-trees is an intuitive and expressive qualitative preference formalism over combinatorial domains, with an edge of lower computational complexity compared to other languages.
- 2 With formulas restricted to attributes as node labels, the language of PLP-trees is highly accurate in modeling preferences arising in practice, and can be effectively learned. Collections of PLP-trees, or PLP-forests, are empirically shown with reduced overfitting and higher accuracy.
- 3 With a further restriction that all attributes occur exact once on every path, the language of LP-trees represent total orders, and thus embrace preference aggregation problems using voting schemas (e.g., positional scoring rules). As the problems are in general NP-hard, we apply ASP and show ASP tools are effective for large instances.

Questions?

Thank you!